

# Distributed Systems - Assignment 1

---

Gowlapalli Rohit - 2021101113

## Question 1

### Distributed BFS

The graph is partitioned into equal parts and each process is assigned a set of vertices. When a vertex is visited, its neighbors are communicated to the corresponding process.

#### Time Complexity:

- **Factors:**
  - Graph traversal (BFS)
  - Data exchange (MPI communication)
- **Graph Partitioning:**  $O(V/\text{size})$  per process
- **BFS Traversal:**
  - $O(V/\text{size} + E)$  per process
  - Communication overhead:  $O(d * V/\text{size})$
  - Total:  $O(V/\text{size} + E + d * V/\text{size})$
- **Gathering Results:**  $O(V)$

**Total Time:**  $O(V/\text{size} + E + d * V/\text{size}) + O(V)$

#### Space Complexity:

- **Per Process:**  $O(V/\text{size} + E/\text{size} + L)$ 
  - Adjacency list
  - Level array
  - Blocked set
- **Total System:**  $O(V + E + L)$

#### Message Complexity:


- **Initial Distribution:**  $O(E)$
- **BFS Traversal:**  $O(V * d)$

**Total Messages:**  $O(V * d)$

#### Summary:

- **Time:**  $O(V/\text{size} + E + d * V/\text{size}) + O(V)$
- **Space:**  $O(V + E + L)$
- **Messages:**  $O(V * d)$

## Scaling Analysis



## Question 2

### Distributed Particle Simulation

The grid is divided into rows and each process is assigned a set of rows. When a particle moves to a different row, it is communicated to the corresponding process.

#### Time Complexity:

- **Initialization:**  $O(K)$ 
  - Reading  $K$  particles and assigning them to processes.
  - Broadcasting  $N, M, K, T$ :  $O(\log P)$
- **Particle Movement & Collision Handling:**  $O(T * K)$ 
  - Per iteration:
    - Updating positions:  $O(K/P)$
    - Communication:  $O(K/P)$
    - Erasing invalid particles:  $O(K/P)$
    - Collision handling:  $O(K/P)$
- **Gathering Results:**

- Sorting particles:  $O(K \log K)$
- Gathering particles:  $O(K + P)$

### Space Complexity:


- **Per Process:**  $O(K/P)$  (particles, collision handling data)
- **Global (Rank 0):**  $O(K)$  (gathered particles)

### Message Complexity:

- **Broadcasting Constants:**  $O(\log P)$
- **Particle Communication:**  $O(T * K)$
- **Gathering Results:**  $O(K + P)$

### Summary:

- Time:  $O(T * K + K \log K)$
- Messages:  $O(T * K + P)$
- Space:  $O(K + K/P)$



## Question 3

- **Nodes and Chunks:**
  - N: Total number of nodes.
  - C: Number of chunks in a file.

- F: Number of failed nodes.
- Replication: Each chunk has a replication factor of 3.
- Storage: The chunk size is `CHUNK_SIZE`.
- **Heartbeat:** All nodes send periodic heartbeat messages to the metadata server.

### Time Complexity

- **Upload:**  $O(C * N)$ 
  - Chunk Upload:  $O(N)$  per chunk (replica node selection)
- **Retrieve:**  $O(C)$ 
  - Replica Search and Retrieval:  $O(1)$  per chunk
- **Search:**  $O(C * W)$ 
  - Node Selection:  $O(C)$
  - Word Search:  $O(W)$  per chunk (where  $W$  is the average words per chunk)
- **Heartbeat:**  $O(N)$ 
  - All nodes send heartbeats to the metadata server.
- **Failover/Recover:**  $O(C * N)$ 
  - Node Selection for Re-replication:  $O(N)$  per chunk

### Space Complexity

- **Metadata Server:**  $O(F * C)$ 
  - File metadata:  $O(F)$
  - Chunk metadata:  $O(C)$  per file
- **Storage Nodes:**  $O(F * C * \text{CHUNK\_SIZE})$ 
  - Chunk data:  $O(C * \text{CHUNK\_SIZE})$  per file

### Message Complexity

- **Upload:**  $O(C)$ 
  - 12 messages per chunk (4 messages per replica)
- **Retrieve:**  $O(C)$ 
  - 4 messages per chunk (metadata request + data retrieval)
- **Search:**  $O(C)$ 
  - 4 messages per chunk (similar to retrieval)
- **Heartbeat:**  $O(N)$ 
  - $N - 1$  heartbeat messages
- **Failover/Recover:**  $O(C)$ 
  - Metadata updates during re-replication

### Summary

- **Overall Time Complexity:**  $O(C * N + C * W + N)$
- **Overall Space Complexity:**  $O(F * C * \text{CHUNK\_SIZE} + F * C)$
- **Overall Message Complexity:**  $O(C * 12 + N)$

## Assumptions

---

- 1D grid partitioning of the graph.

## Question 2

- 1D grid partitioning of the matrix
- It is assumed that particles cannot collide at the start of the simulation.

## Question 3

- Distributed Search is done by fetching only the relevant chunks that have a complete match with search word or if their suffix or prefix is a substring of the search word.
- Heartbeat Interval is 500ms and Failover Interval is 1500ms
- Load balancing is done by always picking storage nodes with the least number of chunks. (When a node fails, the chunks count of the node becomes zero)
- It was assumed that there are no time constraints on the upload and retrieval of chunks. The system is optimized for fault tolerance and load balancing.
- Master Node has maps to store the metadata of files and chunks. It also has a list of all storage nodes.
- Each chunk metadata contains chunk id and the replication nodes.
- Each file metadata also contains offset information for each chunk.

# Distributed File System: Design and Workflow

---

This distributed file system (DFS) leverages MPI (Message Passing Interface) to distribute file operations across multiple nodes, ensuring fault tolerance, scalability, and efficient resource utilization.

---

## Key Components

### 1. File Chunking:

- Files are split into fixed-size chunks (**CHUNK\_SIZE**), enabling parallel storage and retrieval.

### 2. Replication:

- Each chunk is replicated across **REPLICATION\_FACTOR** nodes to provide redundancy.

### 3. Heartbeat Mechanism:

- Nodes send periodic heartbeat messages (**HEARTBEAT\_TAG**) to indicate availability.
- Heartbeats are monitored to detect node failures (**FAILOVER\_INTERVAL**).

### 4. Failover and Recovery:

- Failed nodes are excluded from operations.
- A recovery mechanism re-integrates nodes upon restoration.

### 5. MPI Communication Tags:

- Tags like **UPLOAD\_TAG**, **RETRIEVE\_TAG**, and **SEARCH\_TAG** coordinate operations between nodes.

## 6. Metadata Management:

- `FileMetaData` and `ChunkMetaData` structures store file and chunk information, including replication details.
- 

## Workflow

### 1. File Upload:

- **Master Node** (rank 0): Splits the file into chunks and assigns chunks to nodes using a load-balancing strategy.
- **Worker Nodes**: Store assigned chunks locally.

### 2. File Retrieval:

- Master node queries workers for chunks of a file.
- Worker nodes respond with the requested data.

### 3. Search Operation:

- Searches for a word across file chunks.
- Worker nodes analyze chunks locally and send results to the master.

### 4. Heartbeat Monitoring:

- Master node monitors the heartbeats of all nodes.
- Failure detection triggers failover actions (e.g., halting operations on failed nodes).

### 5. Node Failover and Recovery:

- Nodes can be marked as failed using a `failover` command.
- A `recover` command reintegrates restored nodes into the system.

### 6. System Exit:

- Master sends an `EXIT_TAG` to terminate all worker nodes gracefully.